

Self-Supervised SimCLR Pretraining for Automated Wound Segmentation Using U-Net Architectures

Nadia Jelani

Department of Computing, College of Business, Technology and Engineering
Sheffield Hallam University
Sheffield, United Kingdom
Email: nadiajelani0@gmail.com

Abstract—Accurate wound segmentation is essential for effective chronic wound management, yet clinical practice often relies on subjective visual assessment. This paper presents a comprehensive deep learning framework that integrates self-supervised representation learning (SimCLR) with U-Net for automated wound segmentation. By leveraging contrastive pretraining on 5,000 unlabeled wound images, the system learns robust feature representations that significantly improve downstream segmentation accuracy. We developed a modularized Python package (“woundseg”) with complete pipelines for preprocessing, segmentation, explainability (SHAP, Grad-CAM), and automated clinical reporting. Experimental results on 1,247 annotated wound images demonstrate that SimCLR-pretrained U-Net achieves 5.1% improvement in IoU (0.78→0.82) and 3.4% increase in Dice coefficient (0.87→0.90) compared to baseline U-Net. The framework includes comprehensive explainability analysis, clinical metrics calculation, and automated report generation, demonstrating the potential of self-supervised learning for clinical AI applications in wound care.

Index Terms—Wound Segmentation, U-Net, SimCLR, Self-Supervised Learning, Medical Imaging, Clinical AI

I. INTRODUCTION

Chronic wounds affect over 6.5 million patients annually in the United States alone, with diabetic foot ulcers and pressure sores representing the most prevalent conditions. The economic burden exceeds \$50 billion annually, with wound care costs increasing by 8% each year. Manual wound assessment methods are prone to subjectivity, inter-observer variability, and inconsistent documentation, leading to suboptimal treatment outcomes and delayed healing.

Deep learning approaches, particularly U-Net architectures [1], have demonstrated remarkable success in medical image segmentation tasks. However, the performance of supervised learning methods is fundamentally limited by the availability of large, high-quality annotated datasets. In medical imaging, expert annotation is expensive, time-consuming, and requires specialized clinical knowledge, resulting in limited training data that restricts model generalization and performance.

Self-supervised learning (SSL) has emerged as a powerful paradigm to address data scarcity by leveraging large unlabeled datasets to learn meaningful representations. SimCLR [2] has shown particular promise in learning robust visual representations through contrastive learning, achieving state-of-the-art performance on various downstream tasks. We propose a comprehensive framework that integrates SimCLR

pretraining with U-Net for automated wound segmentation, addressing the critical need for accurate, objective wound assessment tools that can support clinical decision-making and improve patient outcomes.

II. METHODS

A. Dataset

We used a diverse dataset of 1,247 chronic wound images with expert-annotated ground-truth masks, spanning diabetic foot ulcers, pressure sores, and surgical wounds. The dataset includes images from various lighting conditions, wound stages, and patient demographics. Images were preprocessed by resizing to 128×128 pixels and normalized to $[0,1]$ range. Data augmentation included random crops, horizontal/vertical flips, brightness adjustments ($\pm 20\%$), and contrast variations to improve model robustness.

B. SimCLR Pretraining

A ResNet-50 backbone was trained with SimCLR on 5,000 unlabeled wound images using contrastive learning. The contrastive loss function maximizes agreement between differently augmented views of the same image while minimizing agreement with views from different images:

$$\mathcal{L}_{contrastive} = -\log \frac{\exp(\text{sim}(z_i, z_j)/\tau)}{\sum_{k=1}^{2N} \mathbf{1}_{k \neq i} \exp(\text{sim}(z_i, z_k)/\tau)} \quad (1)$$

where z_i and z_j are representations of augmented views, τ is the temperature parameter (0.07), and sim is cosine similarity. Strong augmentations including random cropping (0.08-1.0), color jittering (strength 0.8), and Gaussian blur improved representation learning. The pretrained encoder weights were then used to initialize the U-Net architecture.

C. U-Net Architecture and Fine-Tuning

The U-Net encoder was initialized with SimCLR-pretrained ResNet-50 weights, while the decoder was trained from scratch. The architecture includes skip connections between encoder and decoder layers to preserve fine-grained details. Focal Tversky loss was employed to handle class imbalance:

$$\mathcal{L}_{FT} = (1 - \text{Tversky})^\gamma \quad (2)$$

where $\gamma = 1.5$ and Tversky index is defined as:

$$\text{Tversky} = \frac{TP}{TP + \alpha FN + \beta FP} \quad (3)$$

with $\alpha = 0.3$ and $\beta = 0.7$ to penalize false negatives more heavily. Training used Adam optimizer with initial learning rate of 1×10^{-4} , cosine annealing scheduling, and early stopping with patience of 15 epochs.

D. Modular Pipeline Implementation

We developed “woundseg,” a comprehensive Python package with clear modularity:

- **Preprocessing:** Image normalization, resizing, and augmentation
- **Segmentation:** SimCLR-pretrained U-Net inference
- **Postprocessing:** Morphological operations, connected component analysis
- **Explainability:** SHAP and Grad-CAM visualization
- **Clinical Reporting:** Automated PDF generation with metrics

The package provides a CLI interface for reproducible workflows:

```
ws analyze --image sample.jpg --report --explain
```

E. Evaluation Metrics

We evaluated performance using multiple metrics:

- **IoU (Intersection over Union):** $\frac{|A \cap B|}{|A \cup B|}$
- **Dice Coefficient:** $\frac{2|A \cap B|}{|A| + |B|}$
- **Precision:** $\frac{TP}{TP + FP}$
- **Recall:** $\frac{TP}{TP + FN}$

where A is the ground truth mask, B is the predicted mask, and TP , FP , FN are true positives, false positives, and false negatives respectively.

III. RESULTS

A. Experimental Setup

We evaluated our SimCLR-pretrained U-Net approach on a diverse dataset of chronic wound images. The model was trained using a two-stage approach: first, self-supervised pre-training with SimCLR on unlabeled wound images, followed by supervised fine-tuning on annotated segmentation data.

B. Quantitative Results

Table I shows segmentation performance comparing baseline U-Net and SimCLR-pretrained U-Net across multiple evaluation metrics.

TABLE I
SEGMENTATION PERFORMANCE COMPARISON

Model	IoU	Dice	Precision	Recall
Baseline U-Net	0.78	0.87	0.89	0.85
SimCLR + U-Net	0.82	0.90	0.92	0.88

The SimCLR-pretrained U-Net achieved significant improvements across all metrics, with a 5.1% increase in IoU and 3.4% increase in Dice coefficient compared to the baseline.

Table II provides a comprehensive comparison with state-of-the-art methods in medical image segmentation, demonstrating the competitive performance of our approach.

TABLE II
COMPARISON WITH STATE-OF-THE-ART METHODS

Method	IoU	Dice	Parameters
U-Net (Baseline)	0.78	0.87	31.0M
Attention U-Net	0.79	0.88	34.2M
DeepLabV3+	0.80	0.89	16.7M
SimCLR + U-Net (Ours)	0.82	0.90	31.0M

C. Qualitative Analysis

Figure 1 shows a representative wound image used in our analysis, demonstrating the clinical complexity of chronic wound assessment.



Fig. 1. Original chronic wound image showing irregular boundaries and varying tissue characteristics. The image demonstrates the complexity of automated wound segmentation in clinical settings.

Figure 2 presents the segmentation results, showing the binary mask generated by our SimCLR-pretrained U-Net model.

Figure 3 demonstrates the overlay visualization, combining the original image with the segmentation mask to show the precise wound boundaries.

D. Explainability Analysis

To understand the model’s decision-making process, we generated SHAP (SHapley Additive exPlanations) heatmaps as shown in Figure 4.

The SHAP analysis reveals that our SimCLR-pretrained model focuses on clinically relevant features such as wound edges, tissue texture, and color variations, demonstrating the effectiveness of self-supervised pretraining in learning meaningful representations.

E. Clinical Metrics

Our system provides comprehensive clinical assessment including wound area (7.28 mm²), perimeter (12.4 mm), severity classification (Mild), and healing potential (80%).




Fig. 2. Binary segmentation mask generated by SimCLR-pretrained U-Net. The mask accurately identifies wound boundaries with minimal false positives.



Fig. 3. Wound segmentation overlay showing the original image with predicted wound boundaries highlighted. The overlay demonstrates accurate boundary detection and minimal over-segmentation.

The system was validated against expert clinician assessments on 50 cases, achieving Cohen’s kappa of 0.82 for severity classification. Performance analysis shows 15% improvement in boundary smoothness, 23% reduction in false positive rate, and processing time of 0.8 seconds per image on GPU. Qualitative results indicated smoother wound boundaries and better handling of irregular wound shapes using the SimCLR-pretrained approach.

IV. DISCUSSION

Our results demonstrate that SimCLR pretraining significantly improves wound segmentation performance by learning robust feature representations from unlabeled data. The 5.1% improvement in IoU and 3.4% increase in Dice coefficient indicate that self-supervised learning effectively captures clinically relevant features. The modular “woundseg” package design enhances reproducibility and clinical translation by providing standardized preprocessing pipelines, automated re-




Fig. 4. SHAP explanation heatmap showing the importance of different image regions for segmentation decisions. Red regions indicate high importance for wound classification, while blue regions show low importance.

port generation, and explainable AI features. The two-stage training approach offers data efficiency, improved generalization, and clinical interpretability. Limitations include dataset size constraints and computational requirements for real-time deployment. Future work will focus on temporal analysis, larger datasets, and edge deployment optimization.

V. CONCLUSION

We demonstrated that SimCLR-pretrained U-Net achieves superior segmentation performance over baseline U-Net, with significant improvements in IoU (5.1%) and Dice coefficient (3.4%). The comprehensive “woundseg” package provides a reproducible and extensible framework for clinical AI applications, incorporating explainability features, automated reporting, and clinical validation. Our approach addresses critical challenges in wound care by providing objective, consistent, and accurate assessment tools that can support clinical decision-making and improve patient outcomes.

ACKNOWLEDGMENT

The authors thank clinical collaborators and Sheffield Hallam University for support.

REFERENCES

- [1] O. Ronneberger, P. Fischer, and T. Brox, “U-Net: Convolutional Networks for Biomedical Image Segmentation,” in MICCAI, 2015.
- [2] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, “A Simple Framework for Contrastive Learning of Visual Representations,” in ICML, 2020.
- [3] C. K. Sen, G. M. Gordillo, S. Roy, R. Kirsner, L. Lambert, T. K. Hunt, F. Gottrup, G. C. Gurtner, and M. T. Longaker, “Human skin wounds: a major and snowballing threat to public health and the economy,” *Wound Repair and Regeneration*, vol. 17, no. 6, pp. 763-771, 2009.
- [4] S. R. Nussbaum, M. J. Carter, C. E. Fife, J. DaVanzo, R. Haught, M. Nussgart, and D. Cartwright, “An economic evaluation of the impact, cost, and medicare policy implications of chronic nonhealing wounds,” *Value in Health*, vol. 21, no. 1, pp. 27-32, 2018.